



# AI STACK DOCTOR

v4 · Product Summary

*Know exactly where your AI infrastructure stands — in 90 seconds.*

44

Company Profiles

7

AI Domains

14

Compliance Frameworks

100

Point Score



[ai-stack-doctor.onrender.com](https://ai-stack-doctor.onrender.com)



[github.com/dwnjuguna/AI-Stack-Doctor](https://github.com/dwnjuguna/AI-Stack-Doctor)

Open Source · MIT License · Free Forever · Global by Design

# AI is everywhere. But is it working?

88/130

Marketing leaders NOT using AI  
to manage their own AI stack

8%

Of organizations report full  
confidence in AI governance

73%

Have a formal GenAI policy —  
but a policy is not a finish line

*Source: State of Martech 2026 — Brinker & Riemersma (n=208 marketing & martech leaders)*

*"A policy is a starting point, not a finish line — the gap between 'we have a policy' and 'we have the infrastructure to enforce it' is where most organizations are still working things out."*

# 7 domains. 100-point score. 90 seconds.

**14 pts**

## GenAI / LLMs

Foundation models, fine-tuning,  
RAG

**14 pts**

## Agentic AI

Orchestration, tool use, agents

**14 pts**

## Machine Learning

Training frameworks, model  
lifecycle

**14 pts**

## Data Engineering

Pipelines, feature stores,  
warehouses

**14 pts**

## AI Platforms

Internal ML platforms, model  
serving

**14 pts**

## MLOps / LLMOps

Monitoring, CI/CD for models

**16 pts**

## Cloud AI Services

AWS / GCP / Azure AI maturity



Healthy 80–100



Needs Attention 60–79



At Risk < 60

# What's new in v4



## Industry Intelligence (REWired)

9 sector profiles calibrated to McKinsey's REWired framework — prescriptions weighted by highest-value AI domains per industry



## 44 Company Profiles

US, Europe, Asia, Latin America, Africa — including the first-ever African company AI benchmarks



## Deprecation Risk Intel

Powered by State of Martech 2026 — flags tools at high risk of absorption by ChatGPT, Claude, or Gemini



## Smart Intake Forms

4 persona-tailored forms — Consultant, Executive, Marketer, General — ground truth before the first call



## MCP Server — Live

stdio MCP server ships with the repo. Audit any company directly from Claude Desktop. Pro tier adds Streamable HTTP for any MCP client



## ROI Layer

Every recommendation includes gap cost, fix cost, projected ROI, and payback period across 9 domains



## AI Org Health

New report section: CAIO signals, AI platform team maturity, production depth assessment



## Agentic Scheduler

Autonomous audit scheduling, change detection, and Slack/email alerts. Zero external dependencies

KEY FEATURES

# Built for consultants, executives, and technical teams.



## 44 Company Profiles

5 regions · 28 industries · US, Europe, Asia, Latin America, Africa  
First AI benchmarking tool with African company profiles



## 14 Compliance Frameworks

EU AI Act · GDPR · CCPA · HIPAA · ISO 42001 · NIST AI RMF  
PCI-DSS v4 · SOC 2 · China GenAI Regs · UK AI Regulation + 4 more



## Agentic Scheduler

Autonomous weekly/monthly audits · change-detection alerts  
Slack & email webhooks · zero external dependencies



## ROI Layer

Gap cost · fix cost · projected ROI · payback period  
Every recommendation quantified across 9 domains



## Deprecation Risk

Jasper · Copy.ai · Writesonic · Anyword flagged HIGH risk  
Powered by State of Martech 2026 consolidation data



## Intake Forms

4 persona-tailored forms: Consultant · Executive · Marketer · General  
Intake JSON = high-confidence ground truth for the audit engine

# Four ways to run AI Stack Doctor.

01

## Interactive CLI

```
python3 ai_stack_health_agent_v3.py
```

Best for: Individual analysts, quick audits

02

## REST API

```
python3 ... --api → POST /audit
```

Best for: CI/CD integration, automation

03

## Dashboard

```
python3 dashboard_server.py
```

Best for: Executives, visual trend analysis

04

## MCP / Claude Desktop

```
Ask Claude: "Audit Stripe's AI stack"
```

Best for: AI-native workflows, zero context-switching

# 6-tool agentic loop — every audit, every time.

01

## **detect\_ai\_stack**

8 searches × 7 domains — job postings, blogs, vendor sites, GitHub

02

## **research\_stack\_health**

Deprecations, vendor stability, G2 sentiment, State of Martech 2026 intel

03

## **check\_ai\_integrations**

Pipeline health, data flows, model serving, observability gaps

04

## **audit\_governance\_and\_ownership**

14 compliance frameworks + ROI context + AI Org Health + REWIRED industry weighting

05

## **detect\_redundancies\_and\_gaps**

Capability overlaps, duplicate tools, wasted spend estimation

06

## **benchmark\_against\_peers**

3-company comparison from 44 pre-loaded intelligence profiles

OUTPUT → [Executive Summary](#) · [Stack Inventory](#) · [Deprecation Flags](#) · [7 Domain Scores](#) · [Peer Benchmarks](#) · [ROI Analysis](#) · [PDF Export](#)

# Up and running in 4 commands.

1

Clone

```
git clone https://github.com/dwnjuguna/AI-Stack-Doctor
```

2

Install

```
pip3 install anthropic ddgs rich flask reportlab gunicorn
```

3

API Key

```
export ANTHROPIC_API_KEY="sk-ant-..."
```

4

Run

```
python3 dashboard_server.py
```

